

R bootcamp, Module 3: Calculations

August 2013, UC Berkeley

Chris Paciorek

Vectorized calculations

Note to BB: remember to start recording.

As we've seen, R has many functions that allow you to operate on each element of a vector all at once.

```
vals <- rnorm(1000)
chi2vals <- vals^2
chi2_df1000 <- sum(chi2vals)
# imagine if the code above were a loop, or three separate loops
```

Advantages:

- much faster than looping
- easier to code
- easier to read and understand the code

Sometimes there are surprises in terms of what is fast, as well as tricks for vectorizing things in unexpected ways:

```
vals <- rnorm(1e+06)
system.time(trunc <- ifelse(vals > 0, vals, 0))
```

```
##      user  system elapsed
##    0.16    0.00    0.15
```

```
system.time(vals <- vals * (vals > 0))
```

```
##      user  system elapsed
##    0.02    0.00    0.02
```

```
tmp <- as.character(vote$age60)
vote$ageMin <- substring(tmp, 1, 2)
vote$ageMin[1:5]
```

```
## [1] "18" "18" "30" "30" "30"
```

What am I doing with `vals * (vals > 0)`? What happens behind the scenes in R?

If you use a trick like this, having a comment in your code is a good idea.

Linear algebra

R can do essentially any linear algebra you need. It uses system-level packages called BLAS (basic linear algebra subroutines) and LAPACK (linear algebra package). Note that these calculations will be essentially as fast as if you wrote C code because R just calls C and Fortran routines to do the calculations.

The BLAS that comes with R is fairly slow. It's possible to use a faster BLAS, as well as one that uses multiple cores automatically. This can in some cases give you an order of magnitude speedup if your work involves a lot of matrix manipulations/linear algebra. More details in Module 11.

Vectorized vector/matrix calculations

Recall that `+`, `-`, `*`, `/` do vectorized calculations:

```
A <- matrix(1:9, 3)
B <- matrix(seq(4, 36, by = 4), 3)
```

```
A + B
```

```
##      [,1] [,2] [,3]
## [1,]    5   20   35
## [2,]   10   25   40
## [3,]   15   30   45
```

```
A + B[, 1]
```

```
##      [,1] [,2] [,3]
## [1,]    5    8   11
## [2,]   10   13   16
## [3,]   15   18   21
```

```
A * B
```

```
##      [,1] [,2] [,3]
## [1,]    4   64  196
## [2,]   16  100  256
## [3,]   36  144  324
```

```
A * B[, 1]
```

```
##      [,1] [,2] [,3]
## [1,]    4   16   28
## [2,]   16   40   64
## [3,]   36   72  108
```

Matrix/vector multiplication

```
A %*% B[, 1]
```

```
##      [,1]
## [1,]  120
## [2,]  144
## [3,]  168
```

```
A %*% B
```

```
##      [,1] [,2] [,3]
## [1,]  120  264  408
## [2,]  144  324  504
## [3,]  168  384  600
```

```
identical(t(A) %*% A, crossprod(A))
```

```
## [1] TRUE
```

Some decompositions

```
require(fields)
times <- seq(0, 1, length = 100)
R <- exp(-rdist(times)/0.2) # a correlation matrix
e <- eigen(R)
range(e$values)
```

```
## [1] 0.02525 32.85537
```

```
e$ectors[, 1]
```

```
## [1] 0.05195 0.05449 0.05699 0.05946 0.06190 0.06431 0.06669 0.06903
## [9] 0.07133 0.07360 0.07583 0.07802 0.08017 0.08227 0.08433 0.08635
## [17] 0.08833 0.09025 0.09213 0.09396 0.09575 0.09748 0.09916 0.10079
## [25] 0.10236 0.10388 0.10535 0.10676 0.10812 0.10942 0.11066 0.11185
## [33] 0.11297 0.11404 0.11505 0.11599 0.11688 0.11770 0.11846 0.11916
## [41] 0.11980 0.12038 0.12089 0.12134 0.12172 0.12204 0.12230 0.12249
## [49] 0.12262 0.12269 0.12269 0.12262 0.12249 0.12230 0.12204 0.12172
## [57] 0.12134 0.12089 0.12038 0.11980 0.11916 0.11846 0.11770 0.11688
## [65] 0.11599 0.11505 0.11404 0.11297 0.11185 0.11066 0.10942 0.10812
## [73] 0.10676 0.10535 0.10388 0.10236 0.10079 0.09916 0.09748 0.09575
## [81] 0.09396 0.09213 0.09025 0.08833 0.08635 0.08433 0.08227 0.08017
## [89] 0.07802 0.07583 0.07360 0.07133 0.06903 0.06669 0.06431 0.06190
## [97] 0.05946 0.05699 0.05449 0.05195
```

```
sv <- svd(R)
U <- chol(R)

devs <- rnorm(100)
Rinvb <- solve(R, devs) #  $R^{-1}b$ 
Rinv <- solve(R) #  $R^{-1}$  -- try to avoid this
```

Pre-allocation

This is slow.

```
vals <- 0
n <- 50000
system.time({
  for (i in 1:n) vals <- c(vals, i)
})
```

```
## user system elapsed
## 1.73 0.00 1.72
```

The same holds for using `rbind()`, `cbind()`, or adding to a list, one element at a time.

This is slow and unclear:

```
vals <- 0
n <- 50000
system.time({
  for (i in 1:n) vals[i] <- i
})
```

```
##      user  system elapsed
##    1.84    0.00    1.87
```

Thoughts on why these are so slow? Think about what R might be doing behind the scenes

The answer is to pre-allocate memory

This is not so slow. (Please ignore the for-loop hypocrisy and the fact that I could do this as `vals <- 1:n`.)

```
n <- 50000
system.time({
  vals <- rep(NA, n)
  # alternatively: vals <- as.numeric(NA); length(vals) <- n
  for (i in 1:n) vals[i] <- i
})
```

```
##      user  system elapsed
##    0.08    0.00    0.08
```

Here's how to pre-allocate an empty list:

```
vals <- list()
length(vals) <- n
head(vals)
```

```
## [[1]]
## NULL
##
## [[2]]
## NULL
##
## [[3]]
## NULL
##
```

```
## [[4]]
## NULL
##
## [[5]]
## NULL
##
## [[6]]
## NULL
```

apply

Some functions aren't vectorized, or you may want to use a function on every row or column of a matrix/data frame, every element of a list, etc.

For this we use the `apply()` family of functions.

```
mat <- matrix(rnorm(100 * 1000), nr = 100)
row_min <- apply(mat, MARGIN = 1, FUN = min)
col_max <- apply(mat, MARGIN = 2, FUN = max)
```

There are actually some even faster specialized functions:

```
row_mean <- rowMeans(mat)
col_sum <- colSums(mat)
```

lapply() and sapply()

```
myList <- list(rnorm(3), rnorm(3), rnorm(5))
lapply(myList, min)
```

```
## [[1]]
## [1] -1.753
##
## [[2]]
## [1] -1.799
##
## [[3]]
## [1] -2.312
```

```
sapply(myList, min)
```

```
## [1] -1.753 -1.799 -2.312
```

Note that we don't generally want to use `apply()` on a data frame.

You can use `lapply()` and `sapply()` on regular vectors, such as vectors of indices, which can come in handy, though this is a silly example:

```
sapply(1:10, function(x) x^2)
```

```
## [1] 1 4 9 16 25 36 49 64 81 100
```

Here's a cool trick to pull off a particular element of a list of lists:

```
params <- list(a = list(mn = 7, sd = 3), b = list(mn = 6, sd = 1), c = list(mn = 2,
  sd = 1))
sapply(params, "[", 1)
```

```
## a b c
## 7 6 2
```

Think about why this works.

Hint:

```
test <- list(5, 7, 3)
test[[2]]
```

```
## [1] 7
```

```
# `[`(test, 2) # need it commented or R Markdown processing messes it
# up...
```

```
# `+`(3, 7)
```

And more `apply()`s

There are a bunch of `apply()` variants, as well as parallelized versions of them:

- `tapply()`, `vapply()`, `mapply()`, `rapply()`, `eapply()`
- for parallelized versions see Module 11 or `?clusterApply`)

Tabulation

- Sometimes we need to do some basic checking for the number of observations or types of observations in our dataset
- To do this quickly and easily, `table()` is our friend
- Let's look at our observations by year and grade

```
unique(vote$pres04)
```

```
## [1] 1 2 3 9 NA 0 4
```

```
tbl <- table(vote$race, vote$pres04)
tbl
```

```
##
##           0      1      2      3      4      9
## white    111 26184 33045  417    14   409
## black     18  6183   824    56     0    21
## hispanic/latino 6  2665  1639   34     3    49
## asian      0   626   384     7     1     2
## other     16  1036   653    22     0    32
```

```
round(prop.table(tbl, margin = 1), 3)
```

```
##
##           0      1      2      3      4      9
## white    0.002 0.435 0.549 0.007 0.000 0.007
## black    0.003 0.871 0.116 0.008 0.000 0.003
## hispanic/latino 0.001 0.606 0.373 0.008 0.001 0.011
## asian    0.000 0.614 0.376 0.007 0.001 0.002
## other    0.009 0.589 0.371 0.013 0.000 0.018
```

```
table(vote$race, vote$pres04, vote$sex)
```

```
## , , = male
```

```
##
```

```
##
```

```
##
```

```
##           0      1      2      3      4      9
## white    55 11200 15582  230     9   262
## black     6  2440   419    37     0    13
## hispanic/latino 3  1140   784    18     3    28
## asian     0   300   181     2     1     1
```



```
##      other          7   464   335      9      0   25
##
## , , = female
##
##
##           0      1      2      3      4      9
##  white      55 14922 17387   186      5   147
##  black      11  3688   401    19      0      8
##  hispanic/latino  3  1505   845    15      0   21
##  asian       0   325   201     4      0      1
##  other       9   566   310    12      0      7
```

```
with(vote[vote$sex == "female", ], table(pres04, race))
```

```
##           race
## pres04 white black hispanic/latino asian other
##      0      55      11              3      0      9
##      1 14922  3688          1505   325   566
##      2 17387   401          845   201   310
##      3   186     19          15     4   12
##      4      5      0              0      0      0
##      9   147      8              21     1      7
```

Can you figure out what `with()` does just by example?

Stratified analyses I

Often we want to do individual analyses within subsets or clusters of our data. As a first step, we might want to just split our dataset by a stratifying variable.

```
subsets <- split(earnings, earnings$race)
length(subsets)
```

```
## [1] 5
```

```
subsets[["9"]]
```

```
##           earn height1 height2 sex race hisp ed yearbn height
## 239      NA      NA      NA    2    9    9 16      99      NA
## 794      NA       5       5    2    9    2 16      34      65
## 966       0       5       1    2    9    1 16      67      61
```

```
## 1027    NA      5      5  2   9   2 12   31   65
## 1054      0      5      3  2   9   2  8   38   63
## 1507  4416      5      2  2   9   2  6   18   62
## 1676 18000      5      6  1   9   1 17   64   66
## 1713 53000      5      3  2   9   1 17   58   63
```

The `%in%` operator can also be helpful.

```
sub <- earnings[earnings$race %in% c(1, 2, 4), ]
```

Stratified analyses II

Often we want to do individual analyses within subsets or clusters of our data. R has a variety of tools for this; for now we'll look at `aggregate()` and `by()`. These are wrappers of `tapply()`.

```
aggregate(earnings, by = list(educ = earnings$educ), FUN = median, na.rm = TRUE)
```

```
##      educ  earn height1 height2 sex race hisp ed yearbn height
## 1      2    NA      5      0.0  1   2   2  2    0.0   60.0
## 2      3  1400      5      2.0  2   1   1  3   26.0   62.0
## 3      4  6600      5      7.5  2   1   2  4   21.0   66.0
## 4      5  1200      5      3.0  2   1   2  5   17.0   63.0
## 5      6  4416      5      5.0  1   1   2  6   20.0   67.0
## 6      7  7000      5      5.0  2   1   2  7   22.0   66.0
## 7      8  6250      5      5.0  2   1   2  8   32.0   66.0
## 8      9  7000      5      4.0  2   1   2  9   38.0   65.0
## 9     10  8000      5      4.0  2   1   2 10   35.0   66.0
## 10    11 10000      5      6.0  2   1   2 11   42.5   67.0
## 11    12 13000      5      5.0  2   1   2 12   52.0   66.0
## 12    13 15000      5      4.0  2   1   2 13   53.0   66.0
## 13    14 20000      5      5.0  2   1   2 14   53.0   66.5
## 14    15 17000      5      4.0  2   1   2 15   51.5   66.0
## 15    16 24000      5      5.0  2   1   2 16   55.0   66.0
## 16    17 25000      5      6.0  2   1   2 17   50.0   67.0
## 17    18 30000      5      6.0  1   1   2 18   47.0   68.0
## 18    98    NA      5      6.0  2   2   2 98   18.0   66.0
## 19    99    NA      6      0.0  1   1   2 99   37.0   72.0
```

```
aggregate(earn ~ ed, data = earnings, FUN = median)
```

```
##      ed  earn
```

```
## 1 3 1400
## 2 4 6600
## 3 5 1200
## 4 6 4416
## 5 7 7000
## 6 8 6250
## 7 9 7000
## 8 10 8000
## 9 11 10000
## 10 12 13000
## 11 13 15000
## 12 14 20000
## 13 15 17000
## 14 16 24000
## 15 17 25000
## 16 18 30000
```

```
aggregate(earn ~ ed + hisp, data = earnings, FUN = median)
```

```
##    ed hisp  earn
## 1 3 1 1400
## 2 6 1 7000
## 3 8 1 5200
## 4 9 1 0
## 5 10 1 0
## 6 11 1 15000
## 7 12 1 12000
## 8 13 1 17500
## 9 14 1 17000
## 10 15 1 2500
## 11 16 1 17000
## 12 17 1 27000
## 13 18 1 42500
## 14 4 2 6600
## 15 5 2 1200
## 16 6 2 4416
## 17 7 2 7000
## 18 8 2 6250
## 19 9 2 8500
## 20 10 2 9000
## 21 11 2 10000
## 22 12 2 14000
## 23 13 2 15000
## 24 14 2 20000
## 25 15 2 17500
```

```
## 26 16      2 24500
## 27 17      2 25000
## 28 18      2 30000
```

```
agg <- aggregate(earn ~ ed + hisp, data = earnings, FUN = median)
xtabs(earn ~ ., data = agg)
```

```
##      hisp
## ed      1      2
## 3    1400      0
## 4       0   6600
## 5       0   1200
## 6    7000  4416
## 7       0   7000
## 8    5200  6250
## 9       0   8500
## 10      0   9000
## 11 15000 10000
## 12 12000 14000
## 13 17500 15000
## 14 17000 20000
## 15  2500 17500
## 16 17000 24500
## 17 27000 25000
## 18 42500 30000
```

Notice the ‘long’ vs. ‘wide’ formats. You’ll see more about that sort of thing in Module 6.

Discretization

You may need to discretize a categorical variable, e.g., by education level:

```
earnings$edLevel <- cut(earnings$ed, breaks = c(0, 11, 12, 15, 16, 18, 99))
levels(earnings$edLevel) <- c("no HS diploma", "HS grad", "some college", "college grad",
                             "grad study", "other")
boxplot(earn ~ edLevel, data = earnings)
```

Stratified analyses III

`aggregate()` works fine when the output is univariate, but what about more complicated analyses than computing the median, such as fitting a set of regressions?




Figure 1: plot of chunk unnamed-chunk-19

```
out <- by(earnings, earnings$edLevel, function(x) {
  if (sum(!is.na(x$earn)))
    lm(earn ~ height, data = x) else NA
})
length(out)
```

```
## [1] 6
```

```
summary(out[[5]])
```

```
##
## Call:
## lm(formula = earn ~ height, data = x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -48563 -16907  -5810   6902 157470
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -121555      43803   -2.78  0.00623 **
## height           2344         648    3.62  0.00041 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 29700 on 148 degrees of freedom
## (75 observations deleted due to missingness)
## Multiple R-squared:  0.0812, Adjusted R-squared:  0.075
## F-statistic: 13.1 on 1 and 148 DF, p-value: 0.000409
```

What's the business with the if statement? Why is this good practice?

Sorting

`sort()` applied to a vector does what you expect.

Sorting a matrix or dataframe based on one or more columns is a somewhat manual process, but once you get the hang of it, it's not bad.

```
ord <- order(earnings$earn, earnings$height, decreasing = TRUE)
# ord <- with(earnings, order(earn, height, decreasing = TRUE))
ord[1:5]
```

```
## [1] 1860 583 351 618 314
```

```
(earnings$earn[ord])[c(1:5, 2025:2029)] # parentheses for clarity
```

```
## [1] 200000 175000 170000 148000 125000 NA NA NA NA NA
```

```
earnings_ordered <- earnings[ord, ]
```

You could of course write your own *sort* function that uses `order()`. More in Module 5.

Merging Data

We often need to combine data across multiple data frames, merging on common fields (i.e., *keys*). In database terminology, this is a *join* operation.

Here's an example using the voting data combined with a built-in R dataset on state information. Warning: the state dataset is *very* old; this is just a toy example.

In this case (as often true) we need to do some machinations to get everything prepared for the merge. The *key* we use is the state name.

```
# PRE-PROCESSING a bit of querying indicates the state numbers are in
# alphabetical order, so we can do this:
numToName <- data.frame(stateNum = 1:50, stateName = row.names(state.x77))
voteWithStateNames <- merge(vote, numToName, by.x = "state", by.y = "stateNum",
  all.x = TRUE, all.y = FALSE)
stateInfo <- data.frame(state.x77)
# need the names as column, not as the row names attribute:
stateInfo$name <- row.names(stateInfo)

# ACTUAL DEMONSTRATION
fullVote <- merge(voteWithStateNames, stateInfo[, c("name", "Population", "Income",
  "Illiteracy", "Life.Exp")], by.x = "stateName", by.y = "name", all.x = TRUE,
  all.y = FALSE)
dim(vote)
```

```
## [1] 76205 18
```

```
dim(fullVote)
```

```
## [1] 76205 23
```

```
fullVote[1:2, ]
```

```
##   stateName state pres04   sex race age9   partyid   income
## 1  Alabama    1        2 female white 50-59 republican $30,000-$49,999
## 2  Alabama    1        2  male white 30-39 republican $15,000-$29,999
##      relign8 age60 age65 geocode      sizeplac brnagain attend
## 1 protestant 45-59 50-64      4 city: 10,000 to 49,999    yes  <NA>
## 2 mormon/lds 30-44 30-39      2      rural    yes  <NA>
##   year region y ageMin Population Income Illiteracy Life.Exp
## 1 2004      1 1     45      3615    3624      2.1    69.05
## 2 2004      1 1     30      3615    3624      2.1    69.05
```

What's the deal with the `all.x` and `all.y` ? We can tell R whether we want to keep all of the `x` observations, all the `y` observations, or neither, or both, when there may be rows in either of the datasets that don't match the other dataset.

Breakout

Problem 1

Suppose we have two categorical variables and we conduct a hypothesis test of independence. The chi-square statistic is:

$$\chi^2 = \sum_{i=1}^n \sum_{j=1}^m \frac{(y_{ij} - e_{ij})^2}{e_{ij}},$$

where $e_{ij} = \frac{y_{i.}y_{.j}}{y_{..}}$, with $y_{i.}$ the sum of the values in the i 'th row, $y_{.j}$ the sum of values in the j 'th column, and $y_{..}$ the sum of all the values. Suppose I give you a matrix in R with the y_{ij} values.

You can generate a test matrix as: `y <- matrix(sample(1:10, 12, replace = TRUE), nrow = 3, ncol = 4)`.

Compute the statistic without *any* loops as follows:

1. Assume you have the e matrix. How do you compute the statistic without loops?
2. How can you construct the e matrix? Hint: the numerator of e is just an *outer product* for which the `outer()` function can be used.

Problem 2

For each combination of sex and education level, find the *second* largest value of earnings amongst the people in the group without any looping.